

PROJECT 1

Machine Learning

Synthetic Speech Detection & Attribution



Group No 16

Abdul Hannan Anjum Chaudhry 2023-11-0058

Muhammad Hamza 2023-11-0359

Zeeshan Ashraf 2023-02-0083




Overview



- Introduction
- Literary Review
- Methodology
- Implementation
- Results
- Analysis
- Conclusion

Introduction




- Manipulation of audio, speech and video has become easier
- logical advances in the area of signal processing, machine learning and deep learning
- classification of algorithms used to generate different synthetic audios
- development of a classifier to identify the algorithm used for the generation of a synthetic audio.



Data Set

- 5000 synthetic audio recordings generated from 5 different algorithms
- dataset of 15000 samples of noisy synthetic speech recordings using noise addition, reverberation, filtering, and lossy compression.



Literary Review



Features Extraction


- zero-crossing rate (ZCR),
- harmonic distribution
- Mel-Frequency Cepstral Coefficients
- Constant-Q transform (CQT),

Classification Models

- Logistic Multi Linear Regression
- Gaussian mixture models (GMMs)
- Multilayer perceptron (MLPs))
- Recurrent neural networks (RNNs)
- Kalman filters
- convolutional neural networks (CNNs)

Methodology

- Pre-Processing
- Feature Extraction
- Feature Selection
- Dimensionality Reduction
- Machine Learning Models Selection



Implementation

Feature Extraction

- The feature extraction in our models is done through the librosa library, which is very efficient MFCC, Spectral Centroid, Chromagram

Dimensionality reduction

- Principle Component Analysis(PCA) to reduce the dimensions because it is difficult to handle large dimensions in the ML models

Model Implementation

- SVM, KNN, Naïve Bayes, Neural Network, and Logistic regression model using the Sklearn library and Hyperparameter tuning to find the best possible parameters

Features extraction through librosa

```
mfcc_list = []
chroma_list = []
spectral_centroid=[]

for idx,aud in enumerate(aud_list):
    print(idx)
    signal, sr = librosa.load(aud)
    signal = signal.flatten()
    mfccs = librosa.feature.mfcc(signal, n_mfcc=13,sr=sr) #extracting mfccs
    #now for delta and delta2 mfccs
    delta_mfcc = librosa.feature.delta(mfccs) #the delta features show how the signals vary with time, will be useful for stuff
    delta2_mfcc=librosa.feature.delta(mfccs,order=2)
    final_mfcc = np.concatenate((mfccs,delta_mfcc,delta2_mfcc))
    #scaling
    final_scaled=np.mean(final_mfcc.T,axis=0) #Scaled features,
    mfcc_list.append(final_scaled)
    chroma_cq = librosa.feature.chroma_stft(y=signal, sr=sr, n_fft=4096) #extracting chroma stft
    chroma_cq = np.mean(chroma_cq.T,axis=0)
    chroma_list.append(chroma_cq)
    cent = librosa.feature.spectral_centroid(y=signal, sr=sr) #extracting spectral centroid
    cent = np.mean(cent.T,axis=0)
    spectral_centroid.append(cent)
```


80 20 split after standardization and PCA implementation

```
scaler = preprocessing.StandardScaler().fit(mfcc_clean_train)
mfcc_clean_train_scaled = scaler.transform(mfcc_clean_train)
print(mfcc_clean_train_scaled[0])
mfcc_clean_test_scaled = scaler.transform(mfcc_clean_test)
from sklearn.decomposition import PCA
pca = PCA(n_components = 0.95)
pca.fit(mfcc_clean_train_scaled)
mfcc_train1 = pca.transform(mfcc_clean_train_scaled)
mfcc_test1 = pca.transform(mfcc_clean_test_scaled)
print("After mfcc shape")
print(mfcc_train1[0].shape)
```

Sklearn implementation of a Neural Network

```
] from sklearn.metrics import classification_report
from sklearn.neural_network import MLPClassifier

def NeuralNetwork(X_train,X_test,y_train,y_test):
    parameters = {
        'learning_rate_init': [0.05, 0.01, 0.005, 0.001],
        'hidden_layer_sizes': [4, 8, 12],
        'activation': ["relu","logistic", "tanh"],
        'batch_size':[1000],
        'max_iter':[10000]}
    final3 = GridSearchCV(estimator=MLPClassifier(),param_grid=parameters,scoring='accuracy',cv=5)
    final4=final3.fit(X_train,y_train)
    predict_list= final4.predict(X_test)

    acc_score = accuracy_score(y_test,predict_list)
    report = classification_report(y_test,predict_list)
    return((acc_score,final4.best_params_['learning_rate_init'],final4.best_params_['hidden_layer_sizes'],final4.best_params_['activation']))
```

Results

Clean dataset


	precision	recall	f1-score	support
0	0.99	1.00	1.00	208
1	0.84	0.89	0.86	202
2	0.87	0.78	0.82	170
3	0.97	0.99	0.98	217
4	1.00	1.00	1.00	203
accuracy			0.94	1000
macro avg	0.93	0.93	0.93	1000
weighted avg	0.94	0.94	0.94	1000



K-Nearest Neighbor

- MFCCs and Chromagram
- Manhattan distance
- class 0, class 3 and class 4 have a near perfect precision score.

Accuracy 93.6%.



Results

Clean and Augmented
dataset combined




K-Nearest Neighbor

- MFCCs and Chromagram
- Manhattan distance
- class 0, class 3 and class 4 have a near perfect precision score.

Accuracy 96.8%.

	precision	recall	f1-score	support
0	0.99	1.00	0.99	816
1	0.93	0.94	0.93	842
2	0.94	0.91	0.93	795
3	0.98	1.00	0.99	772
4	1.00	0.99	1.00	775
accuracy			0.97	4000
macro avg	0.97	0.97	0.97	4000
weighted avg	0.97	0.97	0.97	4000



Results

Clean dataset



Logistic Regression


- MFCCs and Chromagram
- L1 penalty

Accuracy 95%.

For MFCCs+Chroma

The classifier used is Logistic Regression with the l1 penalty

	precision	recall	f1-score	support
0	1.00	1.00	1.00	208
1	0.90	0.85	0.87	202
2	0.83	0.89	0.86	170
3	0.99	0.99	0.99	217
4	1.00	1.00	1.00	203
accuracy			0.95	1000
macro avg	0.94	0.94	0.94	1000
weighted avg	0.95	0.95	0.95	1000



Results

Clean and augmented
dataset Combined




Logistic Regression

- MFCCs and Chromagram
- L1 penalty
- elastic net regression

Accuracy 93.075%
L1 ratio 0.6666

	precision	recall	f1-score	support
0	1.00	1.00	1.00	816
1	0.83	0.86	0.85	842
2	0.85	0.82	0.83	795
3	0.99	0.98	0.98	772
4	1.00	1.00	1.00	775
accuracy			0.93	4000
macro avg	0.93	0.93	0.93	4000
weighted avg	0.93	0.93	0.93	4000



Results

Clean dataset



Naïve Bayes


- MFCCs and Chromagram
- Class 0, 1, 2, 3, and 4

Accuracy 88.8%

For MFCCs+Chroma

The classifier used is Naive Bayes with accuracy 0.888

	precision	recall	f1-score	support
0	1.00	0.91	0.96	199
1	0.77	0.86	0.81	219
2	0.76	0.75	0.75	187
3	0.96	0.91	0.93	191
4	0.99	1.00	0.99	204
accuracy			0.89	1000
macro avg	0.89	0.89	0.89	1000
weighted avg	0.89	0.89	0.89	1000



Results

Clean and augmented
dataset Combined



Naïve Bayes


- MFCCs, Spectral Centroid and Chromagram
- Class 0, 1, 2, 3, and 4
- precision scores of class 1 and 2 are relatively poorer.

Accuracy 85.4 %

For MFCCs+Spec+Chroma

The classifier used is Naive Bayes with accuracy 0.854

	precision	recall	f1-score	support
0	0.97	0.93	0.95	771
1	0.68	0.80	0.74	791
2	0.73	0.65	0.68	857
3	0.96	0.92	0.94	785
4	0.97	1.00	0.98	796
accuracy			0.85	4000
macro avg	0.86	0.86	0.86	4000
weighted avg	0.86	0.85	0.85	4000



Results

Clean dataset




Support Vector machines

- MFCCs and Chromagram
- Radial basis function as the activation function
- 0.01 regularization parameter
- precision scores of class 1 and 2 are much better as compared to classes 0, 3 and 4.

Accuracy 95%

	precision	recall	f1-score	support
0	1.00	1.00	1.00	204
1	0.89	0.87	0.88	203
2	0.87	0.90	0.88	194
3	1.00	0.99	0.99	207
4	1.00	1.00	1.00	192
accuracy			0.95	1000
macro avg	0.95	0.95	0.95	1000
weighted avg	0.95	0.95	0.95	1000



Results

Clean dataset




Support Vector machines

- MFCCs and Chromagram
- Radial basis function as the activation function
- 0.01 regularization parameter
- precision scores of class 1 and 2 are much better as compared to classes 0, 3 and 4.

Accuracy 95%

0	1.00	0.99	0.99	792	
1	0.95	0.93	0.94	797	
2	0.93	0.93	0.93	816	
3	1.00	0.97	0.99	811	
4	0.94	1.00	0.97	784	
accuracy				0.96	4000
macro avg				0.96	4000
weighted avg				0.96	4000



Results

Clean dataset



Neural Networks


- MFCCs, and Chromagram
- Relu as the activation function
- 0.001 as the initial learning rate
- Number of hidden layers was 12
- Class 1 low precision of 79%, class 2 precision of 86%.

Accuracy 94%

For MFCCs+Chroma

The classifier used is Neural Network with accuracy 0.94
with the activation function: relu

	precision	recall	f1-score	support
0	0.99	1.00	0.99	807
1	0.79	0.87	0.83	787
2	0.86	0.79	0.83	799
3	0.98	0.97	0.98	845
4	1.00	1.00	1.00	762
accuracy			0.92	4000
macro avg	0.93	0.92	0.92	4000
weighted avg	0.93	0.92	0.92	4000



Results

Clean and Augmented
dataset combined




Neural Networks

- MFCCs and Chromagram
- Relu as the activation function
- 0.005 as the initial learning rate
- Number of hidden layers was 12

Accuracy 94.625 %

	precision	recall	f1-score	support
0	1.00	1.00	1.00	807
1	0.86	0.88	0.87	787
2	0.89	0.86	0.87	799
3	0.99	0.99	0.99	845
4	1.00	1.00	1.00	762
accuracy			0.95	4000
macro avg	0.95	0.95	0.95	4000
weighted avg	0.95	0.95	0.95	4000



Analysis

The First Analysis of Obtained Results

- Classes 1 and 2 suffer from the worst precision scores as compared to other classes

The Second Analysis of Obtained Results







- Highest accuracy we have achieved on the clean data set until now is 95 percent

The Third Analysis of Obtained Results

- A major feature we decided to overlook was Fourier transform, so now we will investigate that feature





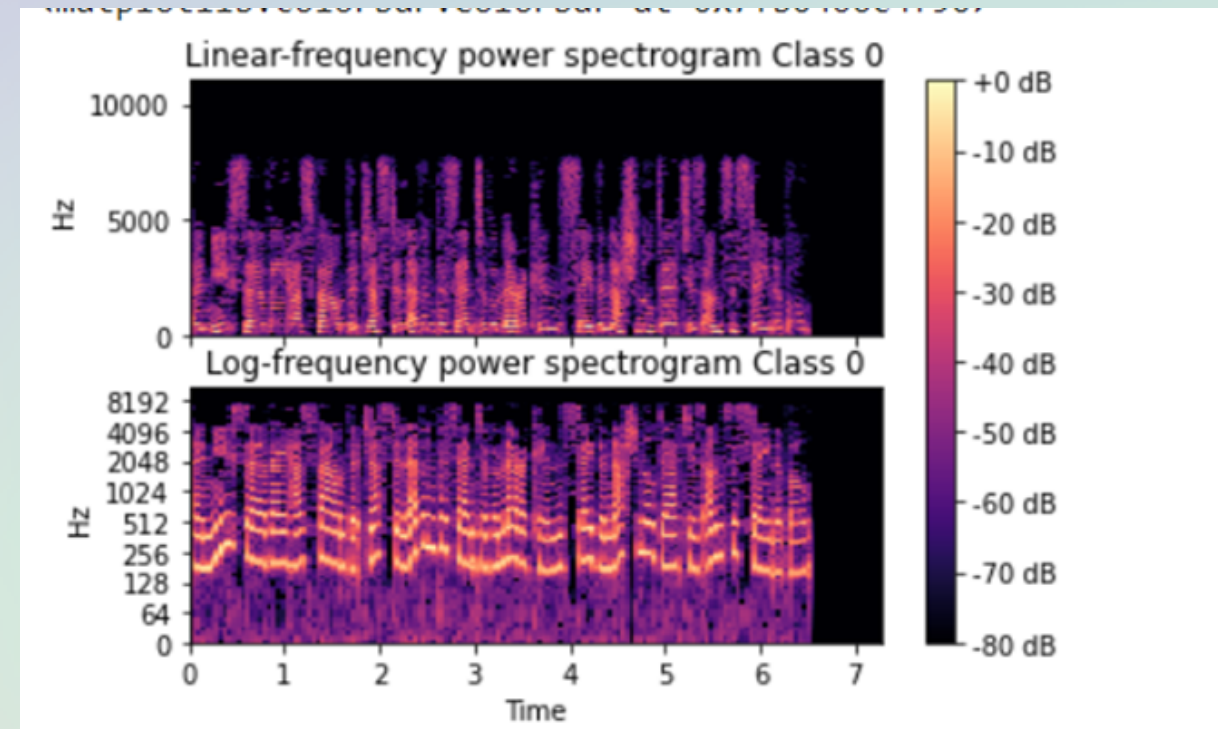
Insights

Amplitude/Signal



Insights

Linear/Log Frequency Power Spectrogram



Conclusion

- Spectral Centroid was not a good feature in terms of synthetic speech algorithm classification
- MFCCs and Chromagram performed the best (in our case)
- Support Vector Machines tend to perform the best with an achieved accuracy of 95% on the clean dataset
- K-Nearest Neighbor performs the best with an achieved accuracy of 96.775% on the augmented + Clean Dataset combined
- Classes 1 and classes 2 suffer from poorer precision scores in KNN algorithm, SVM balances them out a bit with class 4
- Support Vector Machines (which achieved an accuracy of 96.4 on our test data) as the final model



MODEL SELECTION

Support
Vector
Machine

